Nepali Sentiment Analysis of Post-COVID Data

Using XLMRoberta for Text Classification

Amulya Bhandari Sailesh Dahal Sarayu Gautam Tohfa Niraula

April 24, 2025

Department of Computer Engineering
Kathmandu University

Outline

Introduction

Problem Statement

Dataset Description

Data Preprocessing

Tokenization and Encoding

Model Architecture

Training Pipeline

Results

Conclusion

Introduction

What is Sentiment Analysis?

- Sentiment analysis classifies text based on emotion or opinion.
- Categories:
  - Positive — praise, approval
  - Neutral — factual
  - Negative — criticism, disapproval
- Applications:
  - Social media monitoring
  - Product reviews
  - Survey analysis

Problem Statement

What Are We Solving?

- Goal: Classify Nepali-language text into sentiment categories.
- Motivation:
  - Nepali is underrepresented in NLP.
  - Lack of labeled Nepali datasets.
- Objectives:
  1. Clean and preprocess post-COVID Nepali data.
  2. Train a multilingual BERT model.
  3. Evaluate performance using real-world test data.

Dataset Description

About the Dataset

- Source: Nepali COVID/post-COVID text samples.
- Total Samples:
  - Training: 33,602 samples
  - Testing: 8,401 samples
- Labels: 0 = Negative, 1 = Positive, 2 = Neutral
- Common issues:
  - Invalid labels ('o', '-', etc.)
  - Missing values and noisy characters

Data Preprocessing

Data Cleaning Steps

Steps we took:

1. Removed missing and malformed data.
2. Filtered invalid labels.
3. Tokenized using XLM-Roberta tokenizer.
4. Truncated inputs to max length of 256 tokens.

Result: Clean, structured datasets ready for training/testing.

Tokenization and Encoding

Tokenizing with XLM-Roberta

Advantages:

- Supports over 100 languages including Nepali.
- Context-aware encoding using self-attention.
- Subword tokenization handles rare words and typos.

Implementation:

- Used Hugging Face tokenizer from pretrained checkpoint.
- Batch-encoded both train and test sets.

Model Architecture

XLM-Roberta Model Details

Model used: XLM-Roberta-Base

Structure:

- Pretrained encoder: XLM-Roberta
- Classification head: Dense + Softmax layer
- Output: Probabilities over 3 classes (Negative, Positive, Neutral)

Training: PyTorch with mixed precision (autocast enabled)

Training Pipeline

Training Configuration

Training setup:

- Optimizer: AdamW, LR $= 2 \times 10^{-5}$
- Epochs: 10, Batch size: 16
- Loss Function: Cross-entropy
- Platform: Google Colab (GPU)

Libraries used: Hugging Face Transformers, PyTorch, scikit-learn, matplotlib.

Results

Loss and Accuracy Over Epochs

Observations:

- Loss decreased steadily.
- Accuracy reached 93.3

Test Set Evaluation Metrics

Label	Precision	Recall	F1-score
Negative (0)	0.80	0.74	0.77
Positive (1)	0.78	0.83	0.80
Neutral (2)	0.52	0.52	0.52
Overall Accuracy			**74.0%**

Key Insights:

- High precision/recall for Positive/Negative.
- Neutral class more ambiguous $\rightarrow$ lower performance.

Confusion Matrix (Test Set)

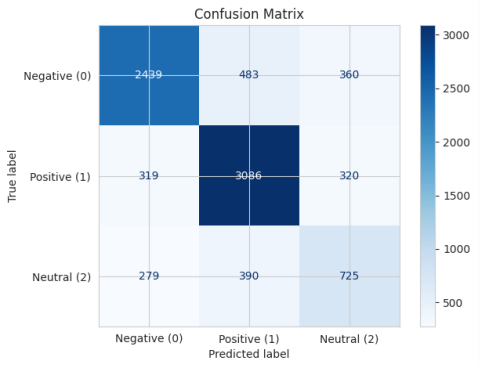

Interpretation: Some overlap between Neutral and other classes $\rightarrow$ expected in real-world data.

Sample Predictions on Unseen Data

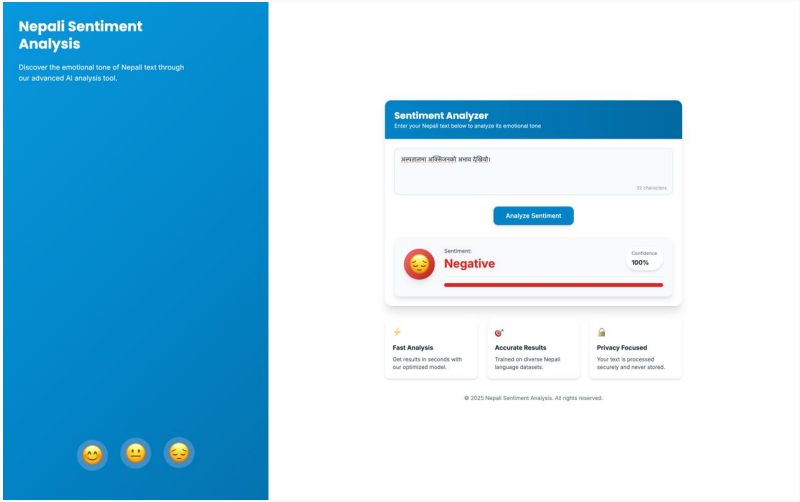

Conclusion

Conclusion and Future Work

Key Takeaways:

- Trained a sentiment classifier on Nepali-language text using XLM-Roberta.
- Achieved 74
- Strong performance on binary sentiment; neutral remains challenging.

Future Improvements:

1. Larger or augmented datasets.
2. Additional validation set for tuning.
3. Model deployment as an API/web service.

Thank You!

Questions or feedback?

Project Resources:

GitHub: `github.com/saileshbro/ai-proj`
We appreciate your time and attention!