Nepali Sentiment Analysis of Post-COVID Data

Using XLMRoberta for Text Classification

Amulya Bhandari   Sailesh Dahal   Sarayu Gautam   Tohfa Niraula

May 2, 2025

Department of Computer Engineering
Kathmandu University

Outline

1

Introduction

What is Sentiment Analysis?

- Sentiment analysis classifies text based on emotion or opinion.
- Categories:
  - Positive — praise, approval
  - Neutral — factual
  - Negative — criticism, disapproval
- Applications:
  - Social media monitoring
  - Product reviews
  - Survey analysis

Problem Statement

What Are We Solving?

- Goal: Classify Nepali-language text into sentiment categories.
- Motivation:
    - Nepali is underrepresented in NLP.
    - Lack of labeled Nepali datasets.
- Objectives:
    1. Clean and preprocess post-COVID Nepali data.
    2. Train a multilingual BERT model.
    3. Evaluate performance using real-world test data.

Dataset Description

About the Dataset

- Source: Nepali COVID/post-COVID text samples.
- Total Samples:
    - Training: 33,602 samples
    - Testing: 8,401 samples
- Labels: $0 =$ Negative, $1 =$ Positive, $2 =$ Neutral
- Common issues:
    - Invalid labels ('o', '-', etc.)
    - Missing values and noisy characters

Data Preprocessing

Data Cleaning Steps

Steps we took:

1. Removed missing and malformed data.
2. Filtered invalid labels.
3. Tokenized using XLM-Roberta tokenizer.
4. Truncated inputs to max length of 256 tokens.

Result: Clean, structured datasets ready for training/testing.

Tokenization and Encoding

Tokenizing with XLM-Roberta

- XLM-Roberta tokenizer is multilingual and supports Nepali language.
- Key features of `tokenizer.batch_encode_plus`:
    - Automatically handles padding and truncation for consistent input length.
    - Generates:
        - `input_ids`: Numerical representation of tokens.
        - `attention_mask`: Identifies non-padded tokens for focus.
- Attention masks ensure the model processes only relevant tokens. It tells the model which sentence is real and which is padding.

Understanding Tokenization and Encoding

- Tokenizing is like chopping a sentence into Lego blocks that the model knows.
- Encoding is turning those blocks into numbers so the model can do math with them.
- Why It Matters:
    - The model sees only numbers.
    - Tokenization keeps input size fixed (e.g., max_length=256). Long texts are trimmed, short ones are padded.

Tokenization Example

Step	Output
Input Sentence	सरकारले अस्पतालमा निःशुल्क उपचार उपलब्ध गरायो।
Tokens	['__सरकारले', '__अस्पतालमा', '__नि', 'ः', 'शुल्क', '__उपचार', '__उपलब्ध', '__गरायो', '।']
Token IDs	[25033, 149007, 946, 6, 153794, 26409, 38071, 87995, 4]

- The model doesn't understand text. It only understands numbers.
- Tokenizer breaks the sentence into subwords it knows (like 'नि', 'ः', 'शुल्क' for 'निःशुल्क').
- The tokenizer assigns a unique ID to each known token.
- These token IDs are what get passed to the model as input.

Model Architecture

XLM-Roberta Model Details

Transformers process all words at the same time (not one-by-one like RNN).

Structure:

- XLM-R supports multiple languages, including Nepali, out of the box.
- Fine-tuning was performed to adapt the model's top layers specifically for sentiment classification.
- Produces probabilities for three sentiment classes: Negative, Positive, and Neutral.

Trained on: PyTorch with mixed precision (autocast enabled)

Model Usage Example

"यो उत्पादन उत्कृष्ट छ।"

Tokenized Input:

- Tokens: [CLS] यो उत्पादन उत्कृष्ट छ । [SEP]
- Converted to input_ids: [101, 2345, 5678, 9101, 1123, 102] (Note: These are example token IDs.)

Model Output (Softmax Probabilities):

- Negative: 0.05
- Neutral : 0.10
- Positive: 0.85

Prediction: → Positive sentiment

- [CLS] and [SEP] are special tokens used by BERT-like models (e.g., XLM-Roberta).
- input_ids are the numerical representation of tokens.
- The model outputs a probability score for each class.
- The highest score (0.85 for Positive) becomes the final prediction.

Training Pipeline
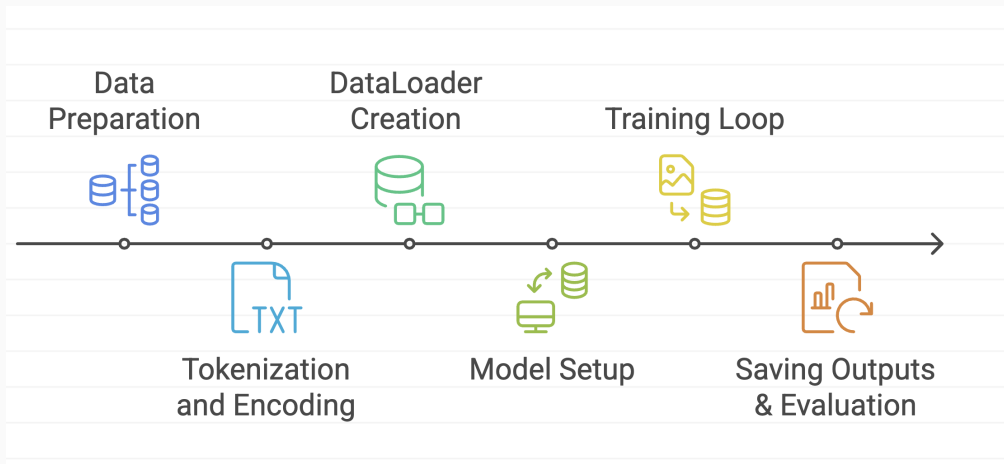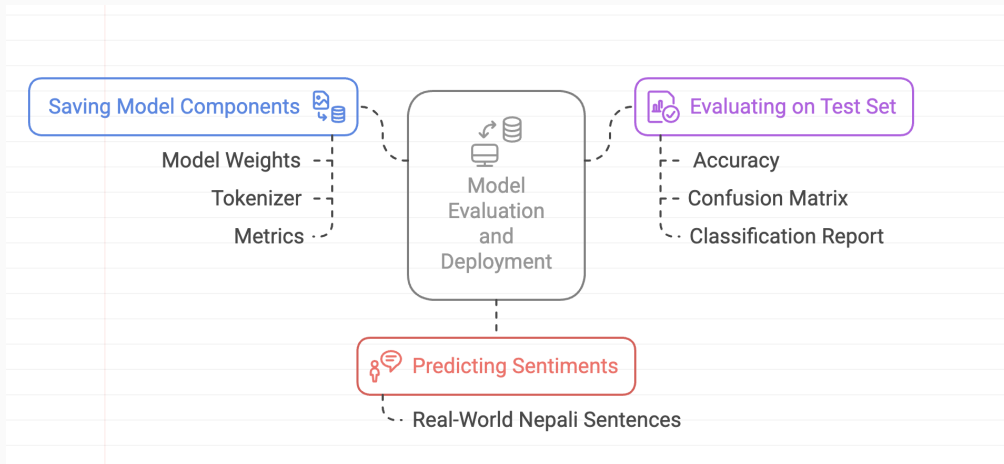
Results

Loss and Accuracy Over Epochs

Training Metrics

Test Set Evaluation Metrics

Label	Precision	Recall	F1-score
Negative (0)	0.80	0.74	0.77
Positive (1)	0.78	0.83	0.80
Neutral (2)	0.52	0.52	0.52
Overall Accuracy			**74.0%**

Key Insights:

- High precision/recall for Positive/Negative.
- Neutral class more ambiguous $\rightarrow$ lower performance.

Confusion Matrix (Test Set)

Conclusion

Conclusion and Future Work

Key Takeaways:

- Trained a sentiment classifier on Nepali-language text using XLM-Roberta.
- Achieved 74% of overall accuracy.
- Strong performance on binary sentiment; neutral remains challenging.

Future Improvements:

1. Larger or augmented datasets.
2. Additional validation set for tuning.
3. Model deployment as an API/web service.

Thank You!

Questions or feedback?

Project Resources:

GitHub: `github.com/saileshbro/ai-proj`

We appreciate your time and attention!